

HIVER AI

AI Support Agent

An AI teammate for fast, safe customer support

KEYNOTE EDITION

Customer message • Context • Answer • Human help



The Experience

A customer message moves from understanding to action.



A production AI teammate, not a basic chatbot.

It reads the request, finds the right context, and responds or hands the case to a person.



Data to Intelligence

Past conversations become usable brand knowledge.

3M+

conversations

24/7

availability

1


shared knowledge layer

Solved cases, product details, and support patterns give the AI the context needed for a useful response.



AI Pipeline

A clear path from incoming message to support decision.



The pipeline classifies intent, retrieves evidence, generates a response, and selects the safest next action.

Knowledge Retrieval Engine

Past solved conversations form the AI memory layer.



Retrieve the most relevant cases

The agent searches for related issues, approved answers, and supporting context before it writes.

RELEVANT EVIDENCE

BRAND CONTEXT

Response Generation

Evidence guides every answer.

01

Retrieve similar cases

Use previous resolutions and policy context.

02

Generate a clear answer

Keep the reply concise, accurate, and on-brand.

03

Verify safety

Check risk before the message reaches a customer.

Human Escalation

The AI knows when a person should take over.

Answer automatically

High confidence, low risk,
clear policy

SAFE TO SEND

Decision based on confidence and risk

Request human support

Uncertain, sensitive, or high-
impact issue

ESCALATE WITH CONTEXT

Evaluation System

Production performance needs measurable evidence.

F1

Intent quality

Reply quality

Customer-ready answers

Human agreement

Safe escalation

Safety

Policy adherence

Failure Radar

Transparent AI surfaces uncertainty and risk.



Knows uncertainty

Signals when the evidence is weak



Detects risk

Flags sensitive and policy-critical cases



Improves continuously

Learns from feedback and corrections



Trust Framework

Accuracy alone does not create confidence.

Trust

=

Quality

+

Safety

+

Human alignment

A trusted support agent combines useful answers with safe behavior and the right human judgement.

Future Roadmap

A system that improves with every useful interaction.

01

Memory

Keep approved
knowledge current

02

Feedback learning

Use human corrections
to improve

03

Deployment

Scale with team
confidence

04

Continuous improvement

Monitor and refine
performance

Built for long-term support quality